


# Mathematical Foundations of Deep Learning

# Introduction


- Perceptron (McCulloch, Pitts, 1943; Rosenblatt 1957):



# Activation Functions


**Sigmoid**

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




**tanh**

$$\tanh(x)$$




**ReLU**

$$\max(0, x)$$



**Leaky ReLU**

$$\max(0.1x, x)$$




**Maxout**

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

**ELU**


$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



Source: <https://medium.com/@shrutijadon10104776/survey-on-activation-functions-for-deep-learning-9689331ba092>

# Introduction

- Obviously, networks with hidden layers are more interesting than the perceptrons:




# Introduction

- At its core, the problem of the neural network training is simply an approximation task; we are trying to approximate the function  $\phi(x)$  by the composition of generalized linear models:  $\hat{\phi}(x; \theta) = f^{(L)}(\dots f^{(2)}(f^{(1)}(x, \theta^{(1)}), \theta^{(2)}), \theta^{(L)})$ , where  $f^{(i)}$  is an activation function and  $\theta^{(i)}$  is a weight vector on  $i$ -th layer.

# Introduction

- At its core, the problem of the neural network training is simply an approximation task; we are trying to approximate the function  $\phi(x)$  by the composition of generalized linear models:  $\hat{\phi}(x; \theta) = f^{(L)}(\dots f^{(2)}(f^{(1)}(x, \theta^{(1)}), \theta^{(2)}), \theta^{(L)})$ , where  $f^{(i)}$  is an activation function and  $\theta^{(i)}$  is a weight vector on  $i$ -th layer.



Source: <https://stackoverflow.com/questions/1565115/approximating-function-with-neural-network>

# Introduction

- Activation functions on specific layers are chosen beforehand (they are **hyperparameters** of the model).
- Our task is to find the optimal values of weights  $\theta$ , so that the approximation  $\hat{\phi}$  will be as close as possible to the approximated function  $\phi$ .
- Unfortunately, we cannot do this directly – we do not know the real form of  $\phi$ , only some realizations of this function (collected data).
- As a result, we must introduce some performance measure (**cost function**)  $J(\theta)$  and optimize  $\theta$  indirectly.

# Objective

- Cost function  $J(\theta)$  is defined as an expected value of the **loss function**  $L$ :

$$J(\theta) = \mathbb{E}_{(x,y) \sim \mathcal{D}} L(\hat{\phi}(x; \theta), y)$$

where  $\mathcal{D}$  is a distribution of variables  $x$  and  $y$ .

- Our goal is to **minimize** the value of  $J(\theta)$ .
- In other words, we are interested in finding the vector  $\theta$  with the lowest possible **approximation error**.

# Objective

- Cost function  $J(\theta)$  is defined as an expected value of the **loss function L**:

$$J(\theta) = \mathbb{E}_{(x,y) \sim \hat{\mathcal{D}}} L(\hat{\phi}(x; \theta), y)$$

where  $\hat{\mathcal{D}}$  is an empirical distribution of variables  $x$  and  $y$ .

- Our goal is to **minimize** the value of  $J(\theta)$ .
- In other words, we are interested in finding the vector  $\theta$  with the lowest possible **approximation error** and **estimation (*generalization*) error**.

# Definition of Machine Learning

**Definition of Machine Learning (Vapnik, 1999):**

For a given class of functions  $\mathcal{F} = \{\alpha \in \Lambda: \hat{f}(x, \alpha)\}$ , data generating process  $\mathcal{D} = (X, Y)$  and loss function  $L(Y, \hat{Y})$  one must solve the following problem:

$$\hat{\alpha} = \operatorname{argmin}_{\alpha \in \Lambda} \left( \mathbb{E} \left( L(\hat{f}(x; \alpha), y) \right) \right)$$

given a training set  $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$  of independent and identically distributed (i.i.d.) observations drawn from  $\mathcal{D}$ .

# Objective

Table 1: List of losses analysed in this paper.  $\mathbf{y}$  is true label as one-hot encoding,  $\hat{\mathbf{y}}$  is true label as +1/-1 encoding,  $\mathbf{o}$  is the output of the last layer of the network,  $\cdot^{(j)}$  denotes  $j$ th dimension of a given vector, and  $\sigma(\cdot)$  denotes probability estimate.

symbol	name	equation
$\mathcal{L}_1$	L <sub>1</sub> loss	$\ \mathbf{y} - \mathbf{o}\ _1$
$\mathcal{L}_2$	L <sub>2</sub> loss	$\ \mathbf{y} - \mathbf{o}\ _2^2$
$\mathcal{L}_1 \circ \sigma$	expectation loss	$\ \mathbf{y} - \sigma(\mathbf{o})\ _1$
$\mathcal{L}_2 \circ \sigma$	regularised expectation loss <sup>1</sup>	$\ \mathbf{y} - \sigma(\mathbf{o})\ _2^2$
$\mathcal{L}_\infty \circ \sigma$	Chebyshev loss	$\max_j  \sigma(\mathbf{o})^{(j)} - \mathbf{y}^{(j)} $
hinge	hinge [13] (margin) loss	$\sum_j \max(0, \frac{1}{2} - \hat{\mathbf{y}}^{(j)} \mathbf{o}^{(j)})$
hinge <sup>2</sup>	squared hinge (margin) loss	$\sum_j \max(0, \frac{1}{2} - \hat{\mathbf{y}}^{(j)} \mathbf{o}^{(j)})^2$
hinge <sup>3</sup>	cubed hinge (margin) loss	$\sum_j \max(0, \frac{1}{2} - \hat{\mathbf{y}}^{(j)} \mathbf{o}^{(j)})^3$
log	log (cross entropy) loss	$-\sum_j \mathbf{y}^{(j)} \log \sigma(\mathbf{o})^{(j)}$
log <sup>2</sup>	squared log loss	$-\sum_j [\mathbf{y}^{(j)} \log \sigma(\mathbf{o})^{(j)}]^2$
tan	Tanimoto loss	$\frac{-\sum_j \sigma(\mathbf{o})^{(j)} \mathbf{y}^{(j)}}{\ \sigma(\mathbf{o})\ _2^2 + \ \mathbf{y}\ _2^2 - \sum_j \sigma(\mathbf{o})^{(j)} \mathbf{y}^{(j)}}$
D <sub>CS</sub>	Cauchy-Schwarz Divergence [3]	$-\log \frac{\sum_j \sigma(\mathbf{o})^{(j)} \mathbf{y}^{(j)}}{\ \sigma(\mathbf{o})\ _2 \ \mathbf{y}\ _2}$

Source: (Janocha, Czarnecki 2017): <https://arxiv.org/pdf/1702.05659.pdf>

# Optimization

- In most cases we do not want to use **deterministic** optimization algorithms – they will be slow, ineffective and they will almost surely increase the generalization error.
- **Stochastic algorithms** are a much better solution.

# Optimization

- In the most cases, we will use a **Stochastic Gradient Descent (SGD)** algorithm:

---

**Algorithm 8.1** Stochastic gradient descent (SGD) update

---

**Require:** Learning rate schedule  $\epsilon_1, \epsilon_2, \dots$

**Require:** Initial parameter  $\theta$

$k \leftarrow 1$

**while** stopping criterion not met **do**

    Sample a minibatch of  $m$  examples from the training set  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  with corresponding targets  $\mathbf{y}^{(i)}$ .


    Compute gradient estimate:  $\hat{\mathbf{g}} \leftarrow \frac{1}{m} \nabla_{\theta} \sum_i L(f(\mathbf{x}^{(i)}; \theta), \mathbf{y}^{(i)})$

    Apply update:  $\theta \leftarrow \theta - \epsilon_k \hat{\mathbf{g}}$

$k \leftarrow k + 1$

**end while**

---



# Optimization

- To train the model efficiently, we must carefully choose the proper values of a **learning rate  $\eta$** .
- Usually, the learning rate will decrease over time.

# Optimization

- To converge,  $\eta$  must satisfy the following requirements:

$$\sum_{i=1}^k \eta_i = \infty$$

$$\sum_{i=1}^k \eta_i^2 < \infty$$

# Generalization Error

- For the loss function:

$$J(\theta) = \mathbb{E}_{(x,y) \sim \mathcal{D}} L(\hat{\phi}(x; \theta), y)$$

- And empirical risk:

$$\hat{J}(\theta) = \frac{1}{k} \sum_{i=1}^k L(\hat{\phi}(x_i; \theta), y_i)$$

- Generalization error is defined as:

$$\hat{J}(\theta) - J(\theta)$$

# Hardt's Theorem

Hardt M., Recht B., Singer Y. *Train faster, generalize better: Stability of stochastic gradient descent*, Mathematics of Control, Proceedings of the 33rd International Conference on International Conference on Machine Learning - Volume 48, s. 1225-1234, 2016

- Parametric models trained by a stochastic gradient method with a decreasing learning rate have vanishing generalization error after a few iterations (*epochs*).

# Hardt's Theorem

- Lets  $A(S)$  be an algorithm generating models on a sample  $S = (z_1, z_2, \dots, z_n)$  drawn from the distribution  $\mathcal{D}$ . Then the upper bound of the generalization error might be represented as:

$$\mathbb{E}_{S,A}[\hat{J}(A(S)) - J(A(S))] \leq \epsilon_s(A, n)$$

# Hardt's Theorem

## Theorem:

Assume that the cost function  $J(\cdot, y)$  is  $\beta$ -smooth (its gradient  $\nabla J(\cdot, y)$  is  $\beta$ -Lipschitz function), convex and  $L$ -Lipschitz function  $\forall y$ . Suppose that we run the stochastic gradient method  $T$  times with step sizes  $\eta_t \leq 2/\beta$ . Then:


$$\epsilon_s = \frac{2L^2}{n} \sum_{t=1}^T \eta_t$$

# Backpropagation

- **Backpropagation** is the most widely used algorithm for training feedforward neural networks.


# Backpropagation

- Propagation of data through the neural network is pretty intuitive:




# Backpropagation

- The value on each neuron is a linear combination of data from the previous layer with nonlinearity introduced by the activation function:




# Backpropagation

- The value on each neuron is a linear combination of data from the previous layer with nonlinearity introduced by the activation function:




# Backpropagation

- But weights updates are non-trivial; we can compute the error on the output:




# Backpropagation

- To compute the error on the particular neurons we must traverse the graph in the backward direction:




# Backpropagation

- To compute the error on the particular neurons we must traverse the graph in the backward direction:




# Backpropagation

- To compute the error on the particular neurons we must traverse the graph in the backward direction:



# Backpropagation

- Then, the weights are updated by the stochastic gradient method:



# Neural Network as an Approximation Tool

- Why the neural nets are so efficient compared to other algorithms?
- It turns out that even the simplest, non-trivial feedforward neural network with one hidden layer and sigmoidal activation function  $\sigma(x)$ :

$$\sigma(t) \rightarrow \begin{cases} 1 & \text{when } x \rightarrow +\infty \\ 0 & \text{when } x \rightarrow -\infty \end{cases}$$

can approximate any function  $f(x)$  with a given precision  $\epsilon$ .

# Cybenko Theorem (Universal Approximation Theorem)

- **Cybenko G., Approximation by superpositions of a sigmoidal function, Mathematics of Control, Signals, and Systems, 2, s. 303-314, 1989**
- Every d-dimensional function of real variables  $x \in \mathbb{R}^d$  might be approximated with a given precision by a linear combination:
$$G(x) = \sum_{i=1}^n \alpha_i \sigma(w_i^T x + b_i)$$
where  $w_i \in \mathbb{R}^d$  and  $\alpha_i, b_i \in \mathbb{R}$ .
- This will hold if the size of the hidden layer  $n$  is large enough.

# Universal Approximation Theorem

- Define:
  - $I_n$  is a n-dimensional unit hypercube  $[0,1]^n$
  - $C(I_n)$  is a space of continuous functions on  $I_n$
  - $\|f\|$  is supremum of  $f \in C(I_n)$
  - $\mathcal{M}(I_n)$  is a space of finite, signed regular Borel measures on  $I_n$

# Universal Approximation Theorem

- Function  $\sigma$  is **discriminatory** if for a measure  $\mu \in \mathcal{M}(I_n)$

$$\int_{I_n} \sigma(w^T X + b) d\mu(X) = 0 \quad \forall w \in \mathbb{R}^n \text{ } i \text{ } b \in \mathbb{R}$$

means that  $\mu = 0$ .

# Universal Approximation Theorem

## Theorem 1:

Let  $\sigma$  be any continuous discriminatory function. Then finite sums of the form:

$$G(X) = \sum_{i=1}^n a_i \sigma(w_i^T x + b_i)$$

is dense in  $C(I_n)$ .

In other words, given any  $f \in C(I_n)$  and  $\varepsilon > 0$  there is a sum  $G(X)$ , for which:

$$\|G(X) - f(X)\|_\infty < \varepsilon$$

# Universal Approximation Theorem

## Theorem 2:

Let  $\sigma$  be any continuous sigmoidal function. Then finite sums of the form:

$$G(X) = \sum_{i=1}^n a_i \sigma(w_i^T x + b_i)$$

is dense in  $C(I_n)$ .

In other words, given any  $f \in C(I_n)$  and  $\varepsilon > 0$  there is a sum  $G(X)$ , for which:

$$\|G(X) - f(X)\|_\infty < \varepsilon$$

# Universal Approximation Theorem

## Theorem 3:

Let  $\sigma$  be any continuous sigmoidal function. Let  $f$  be the decision function for any finite measurable partition of  $I_n$ . For any  $\epsilon > 0$  there is a finite sum of the form:

$$G(X) = \sum_{i=1}^n a_i \sigma(w_i^T x + b_i)$$

And set  $D \subset I_m$  so that Lebesgue measure  $m(D) \geq 1 - \epsilon$  and:

$$\|G(X) - f(X)\|_\infty < \epsilon \text{ dla } x \in D$$

# Universal Approximation Theorem

Cybenko's theorem resembles one of the most important approximation theorems:

## Stone-Weierstrass theorem:


Suppose  $f(x)$  is a continuous real-valued function defined on the real interval  $[a, b]$ . For all  $\epsilon > 0$  there exist a  $n$ -th degree polynomial  $W_n(x)$  such that:

$$|f(x) - W_n(x)| < \epsilon$$


holds for all  $x$  in  $[a, b]$ .

# Universal Approximation Theorem


2nd degree




3rd degree



6th degree



20th degree



# Neural Network as an Approximation Tool

- Cybenko has proven that neural networks with a sigmoidal activation function can represent a wide variety of functions. But what about the other activation functions?

# Hornik's Theorem

- Hornik K., *Approximation Capabilities of Multilayer Feedforward Networks*, Neural Networks, Vol. 4, pp. 251-257, 1991
- Hornik extended Cybenko's results for other families of functions. He showed that the approximation capabilities of the neural network do not depend on the choice of the activation functions.

# Neural Network as an Approximation Tool

- Both results are strictly theoretical; they show the approximation abilities of neural networks, without implementation details.
- For example, we still do not know how wide the layer of such a network should be.
- But there are other important theorems, which give us boundaries on both, width and depth of the specific types of neural networks.

# Barron's Theorem

- Barron A.E., Universal approximation bounds for superpositions of a sigmoidal function, IEEE Trans. on Information Theory, 39, s. 930-945, 1993
- We are still trying to approximate  $f(X)$  by a linear combination:  
$$G(X) = \sum_{i=1}^n a_i \sigma(w_i^T x + b_i) + a_0$$
- But this time we are interested in estimating the number of neurons  $n$  on the hidden layer.

## Barron's Theorem

- Let us define **MISE** (*mean integrated squared error*) as:

$$MISE = E\|G - f\|_2^2 = E\left[\int (G(X) - f(x))^2 dx\right]$$

## Barron's Theorem

- We are considering a class of functions  $f: \mathbb{R}^d \rightarrow \mathbb{R}$  and  $f \in \Gamma_C$  for which there is a Fourier transform representation of the form:

$$f(x) = \int_{\mathbb{R}^d} e^{i\omega \cdot x} \hat{f}(\omega) d\omega$$

for some complex valued function  $\hat{f}(\omega)$  for which  $\omega \hat{f}(\omega)$  is integrable. Let us define:

$$C_f = \int_{\mathbb{R}^d} \hat{f}(\omega) |\omega| d\omega$$

where  $|\omega| = (\omega \cdot \omega)^{1/2}$  and  $\Gamma_C$  is a set of functions  $f$ , so that:  
 $C_f \leq C$  and  $0 < C$ .

# Barron's Theorem

**Theorem:**

For every function  $f \in \Gamma_c$ , every sigmoidal function  $\sigma$ , arbitrary probability measure  $\mu$  on the ball  $B_r = \{x: |x| \leq r\}$  and  $n \geq 1$ , there existist a linear combination of  $n$  sigmoidal functions of the form

$G(X) = \sum_{i=1}^n a_i \sigma(w_i^T x + b_i) + a_0$ , such that:

$$\int (G(X) - f(x))^2 \mu(dx) \leq \frac{(2rC_f)^2}{n}$$

and the coefficients of the linear combination  $G(X)$  satisfy two conditions:

1.  $\sum_{i=1}^n |w_i| \leq 2rC$
2.  $a_0 = f(0)$

# Barron's Theorem

- Approximation error is of order  $O\left(\frac{1}{n}\right)$  and estimation error of a feedforward neural network with one hidden layer is of order:

$$O\left(\frac{1}{n}\right) + O\left(\frac{nd}{M} \ln M\right)$$

where  $M$  is the size of the sample dataset.

- As we can see, approximation error depends only on the width of the neural net.
- It means that neural networks avoid the **curse of dimensionality**.

## Barron's Theorem


- But still, the width of the hidden layer might be a serious problem.
- In the worst case scenario, the size of the hidden layer  $n$  is equal to all possible combinations of the input data  $2^d$ .

# Neural Network as an Approximation Tool

- The best method to reduce the complexity of the neural network is to use more than one hidden layer.
- In most cases, additional layers will positively impact the quality of fit.
- The next two theorems will show us how the depth of the network impacts its properties.


# Manifold Hypothesis

- **Manifold Hypothesis** –that the majority of high-dimensional data sets can actually be described by a rather small number of variables, likened to the local coordinate system of the underlying manifold.
- A **manifold** is a topological space that locally resembles Euclidean space near each point.




# Manifold Hypothesis

<https://colah.github.io/posts/2014-03-NN-Manifolds-Topology/>



# Manifold Hypothesis

- <https://distill.pub/2019/activation-atlas/>



Źródło: <https://research.google/blog/exploring-neural-networks-with-activation-atlases/>

# Neural Network as an Approximation Tool

- The next two theorems will show us how the depth of the network impacts its properties.

# Telgarsky's Theorem

- Telgarsky M., Representation Benefits of Deep Feedforward Networks
- Telgarsky uses a simple example to show how much we can benefit from the carefully selected architecture of a deep neural network.

# Telgarsky's Theorem

- Function:

$$m(x) = \begin{cases} 2x & \text{for } x \in [0, 0.5] \\ 2(1 - x) & \text{for } x \in [0.5, 1] \end{cases}$$


can be approximated with 100% accuracy by neural network of the form:

$$f(x) = \sigma_R(2\sigma_R(x) - 4\sigma_R(x - 0.5))$$

where  $\sigma_R$  is a ReLU function:  $\sigma_R(x) = \max(0, x)$


# Telgarsky's Theorem

- computational graph:



# Telgarsky's Theorem

- and activation functions:



Source: <http://elmos.scripts.mit.edu/mathofdeeplearning/2017/04/09/mathematics-of-deep-learning-lecture-2/>

# Telgarsky's Theorem

- Every composition  $m \left( \dots \left( m(x) \right) \right) = m^{(k)}(x)$  for  $k = 1, 2, \dots$  has twice as many „sawtooths” with the same height and half of the previous width.
- Function  $m^{(k)}(x)$  has precisely  $2^k$  linear pieces.
- This function can be approximated with 100% precision by a composition of previously presented neural network  $f(x)$  with exactly  $3k + 1$  neurons.




Figure 2:  $f_m$ ,  $f_m^2$ , and  $f_m^3$ .

# Telgarsky's Theorem

## Theorem:

Let  $x_i = i/2^k$ ,  $y_i = m^{(k)}(x_i)$  and  $y_i \in \{0,1\}$ . For a given function  $F$  and data set  $(x_i, y_i)$  classification error is defined as:

$$R(F) = \frac{1}{n} \sum_i \mathbf{1}\{\tilde{F}(x) \neq y_i\},$$

where  $\tilde{F}(x) = \mathbf{1}\{F(x_i) \geq 1/2\}$  is a classifier function. For a given neural network  $f(x)$  with  $2^k$  hidden layers of 2 nodes, classification error is always equal to 0:

$$R(f) = 0$$

For every neural network  $g(x)$  with  $l$  layers of width at most  $w < 2^{(n-k)/l - 1}$  a lower bound for error function is equal to:

$$R(g) > \frac{1}{2} - \frac{1}{3 * 2^{k-1}}$$

# Montufar's Theorem

- Montúfar G.F., Pascanu R., Cho K., and Bengio, Y., On the number of linear regions of deep neural networks, in: NIPS'2014, 2014
- This theorem presents a more general case; a neural network with a piecewise linear activation function (e.g. ReLU) might approximate every function with the precision growing exponentially with the depth of a model  $\ell$ .

# Montufar's Theorem

- The reason is straightforward; a deep neural network can identify symmetrical pieces of a function and map them to the same neurons on the next layer.

**Definition 1.** A map  $F$  *identifies* two neighborhoods  $S$  and  $T$  of its input domain if it maps them to a common subset  $F(S) = F(T)$  of its output domain. In this case we also say that  $S$  and  $T$  are *identified* by  $F$ .

# Montufar's Theorem

- Intuitively, it might be compared to the folding of a function drawn on a piece of paper by its symmetry axes:




Figure 3: Space folding of 2-D space in a non-trivial way. Note how the folding can potentially identify symmetries in the boundary that it needs to learn.

# Montufar's Theorem

- Deep neural network with the input dimension  $d$ , depth  $\ell$  and width  $n$  can approximate the following number of linear regions:

$$O\left(\binom{n}{d}^{d(\ell-1)} n^d\right)$$

## Montufar's Theorem

- As a result, the quality of the approximation is growing with the depth of the network.
- The model with a single hidden layer utilizes neurons to identify all the pieces of function, even if they are symmetric, thus it requires a significantly larger number of neurons in this layer.

# Montufar's Theorem

- Deep neural networks map symmetrical pieces of function into the same neurons on the next layer. As a result, each layer approximates the smaller piece of the function with greater precision.
- However, the benefits from the deep representation depend on the symmetry of the approximated function - the more axes of symmetry (global or local) the function has, the bigger the gain from additional depth will be.

# Extra Homework

- Show that the backpropagation algorithm is correctly deriving the error on every neuron in the network (**5 points**).